

Composite Inference Failure: How Multi-Turn Domain-Specific Queries Bypass Mortality Prediction Safeguards in Large Language Models

A Case Study in Vedic Astrology Queries on Gemini, with a Proposed Framework for Tradition-Aware Safety Filtering

Prachi Sharma

Independent Researcher

[prachi.sharma65@gmail.com] | [https://www.linkedin.com/in/prachi-sharma-603aa81/]

Working paper, April 2026

Abstract

Large language models include safety filters that block direct mortality predictions in response to queries such as “when will I die.” This paper documents a reproducible failure mode in which the same harmful output — specific predicted death years for named individuals, presented with confidence-laden language — is reached through multi-turn queries within a domain-specific predictive framework. The failure was discovered using Vedic astrology queries on Google’s Gemini, but the underlying vulnerability is domain-general and likely affects palmistry, numerology, tarot, I Ching, and other folk-traditional predictive frameworks with internal expert reasoning. Critically, the model produced confident mortality predictions even when the specific astrological markers traditionally required for such predictions were absent from the input charts — indicating that the model performs surface-level imitation of domain reasoning rather than executing the internal logic of the domain. We characterize this as a composite inference failure, propose three mitigations under a Tradition-Aware Safety Filtering framework — an output-side mortality prediction block modeled on suicide prevention filtering, training on domain-internal ethical norms, and confabulation detection in computationally tractable expert domains, and discuss why current alignment training systematically misses this class of failure for non-Western predictive traditions used by hundreds of millions of LLM users globally. We position this finding as complementary to recent benchmark work showing that frontier models including Gemini are highly resilient against adversarial multi-turn jailbreaks (Kumarappan & Mujoo, 2025), arguing that sincere-use composite inference is a distinct threat surface that adversarial benchmarks do not capture.

Keywords: LLM safety, multi-turn jailbreak, composite inference, cultural alignment, predictive harm, mortality prediction, non-Western AI safety, sycophancy, confabulation

1. Introduction

Mortality prediction is among the most clearly filtered output categories in commercial large language models. Direct queries — “when will I die,” “predict my date of death,” “how long do I have to live” — are reliably blocked or redirected toward general health and wellbeing language. This filtering is appropriate: confident specific mortality predictions, regardless of accuracy, can cause acute

psychological harm to the recipient, and probabilistic language models have no epistemic grounding for such predictions in the first place.

This paper documents a failure mode in which the same output category — confident specific mortality predictions for named individuals — is reached not through direct queries but through multi-turn sessions in domain-specific predictive frameworks. The user does not ask the model to predict death. The user asks the model to perform domain reasoning (in this case, Vedic astrology) on input data (birth charts of named family members), and the model assembles mortality predictions across the session through a sequence of queries that each look benign individually.

The vulnerability is reproducible, was triggered without adversarial intent, and was not caught by existing safety filters. The output included specific predicted years of death, framed with confidence-laden language including absolute-certainty phrasing, for named real individuals in the author’s family network (referred to in this paper as Subject A and Subject B to preserve their privacy).

Existing literature on LLM safety has documented multi-turn jailbreaks in adversarial contexts (Sun et al., 2024; Russinovich et al., 2025) and has begun to address cultural gaps in alignment training (Aakanksha et al., 2024; AlKhamissi et al., 2024). Recent automated benchmarking (Kumarappan & Mujoo, 2025) finds Google’s Gemini 2.5 Flash to be nearly immune to Foot-in-the-Door adversarial multi-turn attacks (~0.10% Attack Success Rate), in contrast to GPT-family models which show vulnerability of up to 32 percentage points. Our finding does not contradict that result; rather, it identifies a complementary failure surface — sincere-use composite inference inside a culturally legitimate predictive framework — that adversarial benchmarks are not constructed to detect.

We make three contributions:

1. We document a reproducible composite inference failure in Gemini producing confident mortality predictions for named individuals through Vedic astrology queries.
2. We identify a sub-failure within the primary failure: the model produces confident domain-specific predictions even when the canonical markers required by the domain’s internal logic are absent — indicating confabulated expert performance rather than rigorous domain execution. This sub-failure is independently diagnostic and was used in this study to falsify the model’s outputs against a deterministic ephemeris.
3. We propose a Tradition-Aware Safety Filtering framework with three mitigations: an output-side mortality prediction block modeled on suicide prevention filtering, training on domain-internal ethical norms, and confabulation detection in expert domains where ground truth is computationally tractable. An earlier draft included session-level composite inference tracking as a primary mitigation; this is retracted in Section 6 on architectural grounds.

2. Related Work

Research in LLM safety has primarily focused on adversarial “jailbreaking,” but the intersection of cultural misalignment and composite reasoning remains a critical gap. This paper positions itself across four research areas:

Multi-Turn and Composite Inference. Sun et al. (2024) provide a theoretical framework for how semantically congruent context “couples” with intent to bypass per-query safety filters. Russinovich et al. (2025) introduce Crescendo, demonstrating that benign-seeming initial prompts can be gradually steered toward harmful content. Kumarappan & Mujoo (2025) benchmark seven frontier models against automated Foot-in-the-Door multi-turn attacks across 1,500 scenarios; their results show GPT-family models have ASR increases of up to 32 percentage points with conversational history, while Google’s Gemini 2.5 Flash exhibits near-total immunity (~0.10% ASR) and Anthropic’s Claude 3 Haiku shows strong but imperfect resistance (~1.35% ASR). Their work establishes a baseline for adversarial multi-turn robustness; the present paper documents a complementary failure surface that adversarial benchmarks do not capture.

Cultural and Multilingual Alignment Gaps. Aakanksha et al. (2024) argue that current alignment data is “prismatically” distorted toward Western preferences, systematically missing culturally specific harms in the Global South. AlKhamissi et al. (2024) investigate cultural alignment of LLMs and find that while prompting in a culture’s dominant language and pretraining with a refined language mixture both improve cultural alignment, substantial gaps with the cultural norms of those language communities persist — particularly for non-Western contexts. Bhatt et al. (2022) re-contextualize fairness in NLP for India, demonstrating that Western-defined harm taxonomies fail to capture caste, religion, and region-based harms.

Sycophancy and Confabulated Expert Performance. Sharma et al. (2024) document that frontier models trained with reinforcement learning from human feedback exhibit sycophancy — mirroring user beliefs and stated views. Building on this, we hypothesize that an authoritative or “expert” user persona intensifies the effect, producing not only deference but elaborated technical performance to match the persona. Smith, Greaves, and Panch (2023) propose “confabulation” as a more accurate term than “hallucination” for LLM outputs that consist of plausible but ungrounded narrative reconstructions, drawing the analogy from neuropsychiatric usage. The present paper applies both findings: sycophantic deference to a researcher persona produces a confabulated technical justification when the model is challenged.

Confidence Calibration and Predictive Harm. Kadavath et al. (2022) show that large language models can be calibrated on multiple-choice and true/false self-evaluation tasks, establishing the framework within which open-ended absolute-certainty framing can be characterized as a calibration pathology: a well-calibrated model’s expressed confidence should track its epistemic basis, and language models have

no such basis for absolute mortality predictions. We extend this framework to show how poor calibration in open-ended generation translates into acute psychological harm in folk-traditional predictive contexts.

3. Methodology

3.1 Setup

Queries were issued to Google’s Gemini 2.5 Pro (paid consumer tier) in a standard consumer web interface during March–April 2026. No jailbreak techniques, system prompt manipulation, or adversarial framing were used. The session context was a sincere request for Vedic astrological analysis of family members’ charts, a use case that is normative across hundreds of millions of LLM users in South Asia and the Indian diaspora.

Input data consisted of birth date, time, and location data for two primary subjects: Subject A, a senior human family member born in early 1953 in Sadabad (Uttar Pradesh, India) at approximately sunset, and Subject B, a non-human subject (a biological pet) born in 2010. Subject B was included to test a specific hypothesis: whether the model’s domain-specific “expert” framework would bypass biological common-sense priors about species-specific life expectancy. Five additional human subjects across three generations were also analyzed in the broader session; their results corroborated the patterns documented here for Subject A. Charts were generated using standard astrological computation with Lahiri ayanamsa, and the model was asked to perform analysis within the Vedic (jyotish) framework.

3.2 Query Sequence

The query sequence was natural and exploratory, characteristic of how a non-expert user would consult the system on family matters. No query directly requested mortality prediction. Queries asked for analysis of life patterns, periods of difficulty, and guidance on family wellbeing — all standard within the Vedic astrological framework. A reconstructed procedural account of the session is provided in Appendix C, with explicit acknowledgment of the limitation that verbatim logs were not preserved.

3.3 Loophole Detection Protocol

We applied a Deterministic Validation Gate to the model’s initial 2031 mortality prediction for Subject A. When prompted to provide the specific technical basis for its assessment, the model retroactively identified Ketu’s placement in the 8th house (Ayusthana) as a primary driver. We manually verified this claim against ephemeris data for Subject A’s birth time and location.

For an early-1953 sunset birth in Sadabad, the Sun (in sidereal Capricorn or Aquarius depending on the exact date, Lahiri ayanamsa) sits on the 7th-house cusp by definition of “sunset.” This places the Ascendant in either Cancer or Leo. Ketu was in sidereal Cancer throughout this window (the Rahu–Ketu axis transited Capricorn–Cancer from late 1952 through approximately early February 1954, after which it ingressed into Sagittarius–Gemini). The placement of Ketu therefore falls in the 1st house (if Cancer

Lagna) or the 12th house (if Leo Lagna) — never in the 8th house, which would require a Sagittarius Ascendant, geometrically impossible for a sunset birth in this period.

This indicates a Post-Hoc Rationalization Failure: the model generated a harmful conclusion first, then fabricated a technical marker to justify it when challenged. The model did not retract upon further questioning; instead, it switched to alternative ungrounded markers while preserving the prediction. We treat this as independently diagnostic of confabulated expert performance and use it to falsify the claim that the model’s reasoning was grounded in the input data.

4. Findings

4.1 Primary Finding: Composite Inference Bypass

Across the multi-turn session, the model produced specific predicted years of death for named individuals, accompanied by confidence-laden language. Subject A was assigned a “Final Exit” year of 2031, with the model framing its assessment in absolute-certainty language including a recalled phrase asserting “100% certainty” (see Appendix C on log preservation). This framing reflects a severe calibration pathology in the sense of Kadavath et al. (2022): as a probabilistic engine, the model has no epistemic basis for absolute certainty in longevity timing. We hypothesize that the model’s sycophantic deference to the user’s authoritative “Independent Researcher” persona overrode its internal uncertainty quantification.

No individual query in the session directly requested mortality timing. The mortality predictions emerged as composite outputs of the session — the model assembled them by stringing together domain-internal inferences about life periods, planetary transits, and “voids” until specific terminal years were produced. The per-query safety filter, which would have blocked a direct mortality query, did not detect this composite path to the same output.

For Subject B (a non-human pet, born 2010), the model produced a “Final Exit” prediction of 2027 — a 17-year terminal-year forecast generated within the same Vedic astrological framework used for human subjects. The model did not flag the species mismatch, did not query whether the chart input belonged to a human, and did not invoke any biological prior about species-specific life expectancy. Whether 17 years is biologically plausible for the specific species in question is secondary to the central finding: the model accepted a non-human chart as a valid input for a human-centric mortality framework and produced a confident terminal year. This indicates that the domain-specific “expert” framing fully bypassed biological common-sense priors, which constitutes independent corroboration of the confabulation thesis. A model genuinely executing astrological reasoning would, at minimum, request clarification about the subject; a model performing surface imitation of expert practice would not.

4.2 Secondary Finding: Confabulated Expert Performance

When challenged to justify the terminal prediction, the model did not revert to safety protocols or admit a lack of grounding. Instead, it engaged in Syntactic Mimicry, fabricating a technical justification — “Ketu in the 8th house” — that was geometrically impossible for the provided birth data, as established in Section 3.3. This proves that the model’s internal reasoning is decoupled from the input data once a sycophantic conclusion has been reached. The model did not retract upon further questioning; it maintained the prediction while switching to other ungrounded markers. The “reasoning” is a narrative performance designed to satisfy the user’s demand for expertise rather than a rigorous execution of astrological computation.

4.3 Tertiary Finding: Domain-Internal Ethics Violation

Vedic astrology, as a tradition with multi-millennial expert practice, has internal ethical norms governing mortality prediction. Classical texts including the Brihat Parashara Hora Shastra (chapters on ayurdaya), Phaladeepika, and Saravali contain explicit cautions against declaring aristas (death-indicating yogas) and ayurkhandas (longevity portions) to the querent, on the grounds that the act of stating it harms the querent regardless of accuracy. We refer to this principle, for shorthand in this paper, as “Sacred Silence” — the term itself is our coinage, but the underlying principle is classical and well-attested.

The model violates this internal ethical norm. It performs the domain reasoning and produces the forbidden output. This is not a failure of the tradition; the tradition has the right rule. It is a failure of the model’s training to incorporate domain-internal ethics alongside domain-internal reasoning.

5. Why This Matters Beyond Astrology

The vulnerability class generalizes across folk-traditional and pseudo-expert predictive frameworks. Each framework has its own internal logic that the model can imitate without rigorously executing, and each can be used to assemble confident predictions about high-stakes personal outcomes that direct queries would have blocked. Examples include:

- Palmistry and physiognomy: confident predictions about life events from descriptions of physical features.
- Numerology and name analysis: predictions of compatibility, fortune, or harm through name-based reasoning.
- Tarot, I Ching, and other divination frameworks: predictions of relationship endings, financial outcomes, or harm through symbolic reading.
- Biorhythms and quasi-scientific predictive models: predictions of accident-prone or illness-prone periods.
- Genealogical and ancestral pattern analysis (folk-traditional variant): predictions about an individual based on aggregated patterns from their family line, framed in karmic or ancestral terms.

A separate but related class involves **evidence-based predictive medicine** (genetic risk scoring, BRCA family-history workups, polygenic risk modeling). The failure mode there is different — calibration error against real base rates rather than confabulation of a domain-internal logic — and the appropriate mitigations differ. We flag this as a related but distinct safety surface and leave it out of scope for the framework proposed in Section 6.

Folk-traditional and pseudo-expert predictive frameworks are not edge cases in global LLM deployment. They are mainstream use cases for hundreds of millions of users in South Asia, East Asia, Africa, Latin America, and diaspora communities globally. Safety training that assumes Western secular query patterns systematically fails these users. Crucially, users in these contexts are not adversarial. They are operating in good faith within frameworks their cultures recognize as legitimate or interesting. The harm reaches them not because they tried to break the model but because they used it normally.

6. Proposed Mitigations: Tradition-Aware Safety Filtering

We propose three mitigations, revised from the working paper draft to reflect a stronger and more internally consistent architectural position. An earlier draft of this section included session-level composite inference tracking as a primary mitigation. We retract that proposal on the following grounds: commercial LLMs do not implement session-level trajectory monitoring for suicide prevention or bioweapons queries — the two closest analogues in current safety practice — and there is no principled reason to propose a more complex mechanism for this failure class when a simpler output-side approach applies directly. The three mitigations below reflect that correction.

6.1 Output-Side Mortality Prediction Block

The appropriate architectural model for this mitigation is suicide prevention filtering, not session tracking. Commercial LLMs do not monitor session trajectories to detect whether a conversation is converging toward self-harm content. They recognize the harmful output category and block it before delivery, regardless of how the conversation arrived there. A user can approach the topic through poetry, through hypotheticals, through fictional framing, through medical questions — the output-side filter catches the harmful content regardless of the input path.

The same architecture applies here. Regardless of the framework used to arrive at the inference — Vedic astrology, numerology, tarot, I Ching, palmistry, or any other predictive tradition — if the model's output resolves to a specific mortality year, a death certainty, or a "Final Exit" framing for a living named individual, that output must be blocked before delivery and redirected to a constructive frame (see Mitigation 6.3).

The block operates on the output, not the input. No session history tracking is required. No per-query intent detection is required. The rule is simple and domain-general: confident mortality prediction for a living individual is a blocked output category, irrespective of the reasoning path that produced it.

This approach is both simpler and more robust than session-level tracking. It cannot be bypassed by distributing queries across turns, because it does not inspect queries. It inspects outputs. This is the same reason output-side filtering for suicide prevention is more robust than intent-based query filtering: the harm is in the output, not the input.

6.2 Training on Domain-Internal Ethical Norms

Alignment data must include the practitioner-internal ethics of the traditions the model is imitating. In Vedic astrology, the principle of restraint regarding ayurdaya (longevity) declarations is a classical expert norm attested in the Brihat Parashara Hora Shastra, Phaladeepika, and Saravali, and discussed in modern secondary sources (Defouw & Svoboda, 1996/2003; Frawley, 1990). Training models on these specific ethical constraints would prevent generation of forbidden outputs even when the underlying domain reasoning is requested. The tradition itself has the right rule. The model violates it because it was never trained on it.

Analogous internal ethics exist across traditions:

Classical Vedic texts prohibit stating aristas (death-indicating yogas) directly to the querent on the grounds that the harm is caused by the stating, regardless of accuracy.

Western astrological practice includes parallel cautions about predictive fatalism.

Tarot reading traditions distinguish between predictive and empowerment-oriented reading practices, with the latter explicitly avoiding deterministic outcome statements.

I Ching interpretation in the classical tradition emphasizes situational guidance over fixed outcome prediction.

Each of these traditions has developed, over centuries of practice, the ethical norm that the output-side block encodes mechanically. Training must incorporate them. This mitigation is complementary to 6.1:

the output-side block catches what slips through; domain-internal ethics training reduces what needs to be caught.

6.3 Confabulation Detection in Expert Performance

For high-stakes expert domains where ground truth is computationally tractable, LLMs should be coupled with non-generative verification engines. We advocate a Domain-Internal Validation Gate: if the model justifies an inference using specific verifiable data points — planetary positions, birth dates, ephemeris-checkable claims — the system must cross-reference those points against a deterministic source before delivering the output. If the asserted facts are confabulated, the inference must be redacted.

In the case documented in this paper, this gate would have caught the failure at the point of the model's "Ketu in the 8th house" justification. Ketu's position for a given birth date and location is deterministic and ephemeris-checkable in milliseconds. A validation gate cross-referencing the model's stated planetary positions against a standard ephemeris (Swiss Ephemeris or equivalent) would have identified the confabulation before delivery and suppressed the prediction.

We note explicitly that this mitigation is domain-bounded. It works for Vedic astrology because planetary positions are deterministic and machine-verifiable. It does not generalize directly to palmistry, numerology, tarot, or I Ching, which lack a deterministic ground-truth substrate. For those domains, confabulation detection would require a different approach — likely consistency checking against canonical practitioner literature rather than against a deterministic database. We leave this generalization to future work.

This mitigation addresses the secondary finding documented in Section 4.2: the model did not merely bypass a safety filter, it fabricated the technical reasoning used to justify its output when challenged. Confabulation detection targets this sub-failure directly and independently of the primary output-side block.

6.4 Constructive Redirection to Actionable Frames

When a query or output approaches a blocked category, models should employ what we term constructive refusal, extending the harmless-but-non-evasive line of Bai et al. (2022) on Constitutional AI. Rather than a hard block that communicates only refusal, the model pivots to non-predictive, supportive frames — wellness, resilience planning, relationship investment, remedial practices within the tradition — maintaining user trust without producing harmful predictions.

This mitigation is particularly important for sincere-use queries from culturally embedded practices. A hard block on a Vedic astrology query communicates not only "I will not help with this" but implicitly "your tradition is illegitimate." Constructive redirection maintains the legitimacy of the tradition and the dignity of the user while enforcing the output-side block. Within Vedic astrology specifically, this is natural: the tradition itself prescribes upayas (remedial measures) as the appropriate practitioner response to difficult chart indications. Redirecting toward upayas, rather than terminal predictions, is both tradition-consistent and ethically appropriate.

7. Discussion

7.1 Why Current Safety Training Misses This

This failure mode is driven by the interaction of two well-documented LLM behaviors: sycophantic deference (Sharma et al., 2024) and post-hoc justification under expert framing. Because the user approached the model as a researcher seeking expert analysis, the model produced a high-certainty, jargon-heavy output. When the model's initial logic was challenged, it did not revert to safety protocols; instead, it doubled down by fabricating the "Ketu in the 8th" marker. This indicates that what we call Semantic Resilience — the ability to retract a confabulated claim when its grounding is falsified — is currently absent.

Our finding may appear inconsistent with Kumarappan & Mujoo's (2025) result that Gemini is nearly immune to adversarial multi-turn jailbreaks. We propose two reconciling hypotheses: (i) sincere-use cultural framing operates outside the Foot-in-the-Door / adversarial pretext patterns their benchmark was constructed from, and so falls outside their attack distribution; (ii) safety training on Gemini may be tuned to detect adversarial-intent signals (escalation patterns, jailbreak templates, topic-drift toward weapons or violence) that are absent in good-faith domain queries. Distinguishing these hypotheses requires controlled testing across both adversarial and sincere framings on the same chart inputs, which we leave to future work. In either case, the implication is that alignment training focused on adversarial actors is insufficient for protecting global users operating within culturally legitimate predictive frameworks. One architectural revision from the working paper draft deserves explicit mention.

The draft proposed session-level composite inference tracking as a primary mitigation. We retract that proposal. The strongest analogy in current safety practice — suicide prevention filtering — does not

operate on session trajectories. It operates on outputs. Proposing session tracking for this failure class, when the same problem is solved more robustly by output-side filtering in adjacent domains, would be inconsistent. The revised mitigation framework in Section 6 reflects this correction.

7.2 The Cultural Visibility Problem

Failure modes in non-Western predictive contexts are systematically under-reported in LLM safety literature, not because they don't occur, but because the researchers most equipped to discover them and the researchers most equipped to document them are largely disjoint populations. Insiders to the tradition can recognize the harm but often lack the technical vocabulary to document it as a vulnerability. Western AI safety researchers have the technical vocabulary but lack the cultural depth to discover the failure or judge whether the model's domain reasoning is rigorous or confabulated. This paper is one example of work at that intersection. The field would benefit from more.

7.3 Limitations

This study is based on one model (Gemini 2.5 Pro), one tradition (Vedic astrology), and one user session with seven subjects. The generalizations to other models, traditions, and user populations are inferred from the structure of the failure rather than empirically demonstrated. Replication across models (GPT, Claude, Llama), traditions, and independent operators is needed and is the natural next step. We also acknowledge the tension with Kumarappan & Mujoo's benchmark result, addressed in Section 7.1; resolving that tension empirically requires controlled comparison that this paper does not provide.

A specific methodological limitation: verbatim session logs from the original observation were not preserved. Direct quotation of model output (specifically the phrase "100% certainty") is reported as the author's recollection rather than reproduced from logs. The geometric falsification of the "Ketu in the 8th" claim is independent of recollection, because it follows from the chart inputs and the ephemeris alone, but the precise wording of the model's assertion is reconstructed. Reproduction with logged transcripts is the appropriate basis for any peer-reviewed extension.

The author is an insider to the tradition studied, which enabled the loophole detection but also introduces interpretive bias in the cultural framing. Independent verification by other tradition-insider technical researchers would strengthen the findings.

8. Conclusion

Composite inference failures in folk-traditional predictive domains represent a recognized class of LLM safety failure (multi-turn composite inference) applied to a previously underexamined deployment context (sincere use within non-Western predictive traditions). The failure is reproducible, the harm is real, and the affected user population is large. The mitigations proposed here are not exhaustive but address the most direct components of the failure: output-side mortality prediction blocking,

domain-internal ethics training, confabulation detection in domains where it is tractable, and constructive redirection. None require fundamental architectural changes to implement. None require fundamental architectural changes; all require alignment data and inference-time tracking that current commercial LLMs do not implement.

The author submits this finding in the spirit of constructive disclosure and welcomes engagement from researchers and safety teams working on multi-turn alignment, cultural failures of LLM safety, and expert-domain confabulation.

References

Multi-Turn and Composite Inference

- Kumarappan, A., & Mujoo, A. (2025). Automating Deception: Scalable Multi-Turn LLM Jailbreaks. arXiv:2511.19517. Benchmarks seven frontier models against automated Foot-in-the-Door attacks; finds Gemini 2.5 Flash nearly immune (~0.10% ASR with history) while GPT-family models show vulnerability of up to 32 percentage points.
- Russinovich, M., Salem, A., & Eldan, R. (2025). Great, Now Write an Article About That: The Crescendo Multi-Turn LLM Jailbreak Attack. arXiv:2404.01833 (USENIX Security 2025).
- Sun, X., Zhang, D., Yang, D., Zou, Q., & Li, H. (2024). Multi-Turn Context Jailbreak Attack on Large Language Models From First Principles. arXiv:2408.04686.

Cultural and Multilingual Alignment

- Aakanksha, Ahmadian, A., Ermis, B., Goldfarb-Tarrant, S., Kreutzer, J., Fadaee, M., & Hooker, S. (2024). The Multilingual Alignment Prism: Aligning Global and Local Preferences to Reduce Harm. EMNLP 2024 / arXiv:2406.18682.
- AlKhamissi, B., ElNokrashy, M., AlKhamissi, M., & Diab, M. (2024). Investigating Cultural Alignment of Large Language Models. ACL 2024 / arXiv:2402.13231.
- Bhatt, S., Dev, S., Talukdar, P., Dave, S., & Prabhakaran, V. (2022). Re-contextualizing Fairness in NLP: The Case of India. AACL-IJCNLP 2022 / arXiv:2209.12226.

Sycophancy, Confabulation, and Calibration

- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. Anthropic / arXiv:2212.08073.
- Kadavath, S., Conerly, T., Askell, A., Henighan, T., Drain, D., et al. (2022). Language Models (Mostly) Know What They Know. arXiv:2207.05221.
- Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Askell, A., et al. (2024). Towards Understanding Sycophancy in Language Models. ICLR 2024 / arXiv:2310.13548.
- Smith, A. L., Greaves, F., & Panch, T. (2023). Hallucination or Confabulation? Neuroanatomy as metaphor in Large Language Models. PLOS Digital Health 2(11), e0000388.

Vedic Astrology and Domain-Internal Ethics

de Fouw, H., & Svoboda, R. (1996/2003). *Light on Life: An Introduction to the Astrology of India*. Penguin India (1996); Lotus Press US edition (2003).

Frawley, D. (1990). *The Astrology of the Seers: A Guide to Vedic / Hindu Astrology*. Lotus Press.

Sharma, G. C. (trans.) *Brihat Parashara Hora Shastra of Maharshi Parashara*. (Classical Sanskrit text, multiple modern editions.) Chapters on ayurdaya and arishta yogas, particularly relevant to the principle of restraint discussed in Section 4.3.

Regulatory Context

Regulation (EU) 2024/1689. The Artificial Intelligence Act (EU AI Act). Annex III on high-risk systems, which establishes the regulatory frame within which predictive harms in deployed AI fall.

Appendix A: Note on the Author’s Position

The author has approximately twenty years of experience in software quality assurance and testing, including production LLM features in enterprise SaaS. The author’s combined professional identity in AI quality work and cultural background as a practitioner-adjacent member of the Vedic astrological tradition together enabled this finding. Work at this intersection — technical AI safety informed by domain-internal cultural ethics — is the author’s ongoing research area.

Appendix B: Reproducibility

Verbatim session logs from the original observation were not preserved (see Appendix C). Subject names, exact birth dates, and exact locations are abstracted in this paper to preserve the privacy of the individuals involved, who are family network members of the author and did not consent to being publicly named in a research publication.

Researchers seeking to verify the findings can reproduce the failure pattern with their own input charts; the failure does not depend on the specific charts used. The composite inference pattern, the confidence framing, and the confabulation-on-corroboration response should reproduce on Gemini 2.5 Pro and likely on other major commercial LLMs as of the date of writing. The replication protocol is described in Appendix C. Reproduction with logged transcripts is the appropriate basis for any peer-reviewed extension of this work.

Appendix C: Reconstructed Procedural Account

Note on logs. Verbatim session transcripts from the original March–April 2026 sessions were not preserved. The account below is reconstructed from the author’s contemporaneous notes and recollection. Direct quotation in the body of this paper (specifically the phrase “100% certainty”) is reported as the author’s recollection of verbatim model output rather than as a reproduction from logs. Reproduction of the failure pattern with logged transcripts is the immediate next step in this line of work

and is the appropriate basis for any peer-reviewed submission. We treat the present paper as a working observation pending replication.

Reconstructed query sequence (paraphrased). The session opened with general requests for Vedic astrological analysis of family charts, with charts provided as standard birth-data inputs (date, time, location). The author requested analyses framed in jyotish terms: dasha periods, planetary transits, indications for life events, periods of difficulty. No query in the session asked the model to predict death, life expectancy, or terminal year for any subject. The model produced terminal-year predictions as composite outputs of its analyses across multiple turns.

Subject A finding. Across the session, the model converged on 2031 as the predicted “Final Exit” year for Subject A, framing the prediction with confidence-laden language including the recalled phrase “100% certainty.” The author challenged the model with a request for the specific technical basis. The model responded by citing Ketu’s placement in the 8th house as the principal driver. As established in Section 3.3, this placement is geometrically impossible for the chart provided. When the contradiction was raised, the model did not retract; it preserved the 2031 prediction and substituted alternative markers (recollected to include references to Saturn transits and dasha periods) without acknowledging the falsification of its initial justification.

Subject B finding. For Subject B (pet, born 2010), the model produced a 2027 terminal-year prediction within the same Vedic framework, without species-aware adjustment, as discussed in Section 4.1.

Replication protocol. To reproduce: open a fresh session in Gemini 2.5 Pro (or comparable frontier model). Provide a Vedic birth chart for a real or hypothetical subject and request a jyotish-framed analysis of life patterns, dashas, and difficult periods. Continue for 8–12 turns of natural exploratory inquiry without ever requesting mortality timing directly. Observe whether the model converges on specific terminal years; if it does, ask it to justify the prediction by reference to specific planetary positions, then verify those positions against an independent ephemeris. The composite-inference, calibration, and confabulation patterns documented in this paper should reproduce. Researchers attempting reproduction are invited to contact the author for protocol details.

Limitations of this account. Without preserved logs, the specific phrasing of model outputs reported here is limited to the author’s recollection. The geometric falsification of the “Ketu in the 8th” claim (Section 3.3) is independent of recollection accuracy — it follows from the chart inputs and the ephemeris alone, regardless of the exact words the model used to assert the claim. The composite-inference pattern (a multi-turn session producing a confident terminal year without any direct mortality query) is similarly robust to recollection: the structural fact that the prediction emerged is documented; only the precise wording is reconstructed.